

assignment_1

Salma Elshahawy

10/5/2020

Contents

Reading the data	1
Data Exploration	2
Preparing the data for ML model	3
Simple Linear regression Model	3
Fit the data into LR	4
Naive Bayes Model	5
KNN ML Model(k=3)	5
KNN ML Model (k=5)	5
Generate an accuracy comparing table	6
Summary	6

```
set.seed(41)
library(skimr)
```

Reading the data

```
df <- data.frame(
  X = as.factor(c(5, 5, 5, 5, 5, 5, 19, 19, 19, 19, 19, 19, 35, 35, 35, 35, 35, 35, 51, 51, 51, 51, 51,
  Y = c("a", "b", "c", "d", "e", "f", "a", "b", "c", "d", "e", "f", "a", "b", "c", "d", "e", "f", "a", "b", "c", "d", "e", "f",
  label = c("BLUE", "BLACK", "BLUE", "BLACK", "BLACK", "BLACK", "BLUE", "BLUE", "BLUE", "BLUE", "BLACK", "BLUE", "B
)
df
```

```
##      X Y label
## 1    5 a  BLUE
## 2    5 b BLACK
## 3    5 c  BLUE
## 4    5 d BLACK
## 5    5 e BLACK
## 6    5 f BLACK
## 7   19 a  BLUE
## 8   19 b  BLUE
```

```
## 9 19 c BLUE
## 10 19 d BLUE
## 11 19 e BLACK
## 12 19 f BLUE
## 13 35 a BLACK
## 14 35 b BLACK
## 15 35 c BLUE
## 16 35 d BLACK
## 17 35 e BLACK
## 18 35 f BLACK
## 19 51 a BLACK
## 20 51 b BLACK
## 21 51 c BLUE
## 22 51 d BLACK
## 23 51 e BLACK
## 24 51 f BLACK
## 25 55 a BLACK
## 26 55 b BLACK
## 27 55 c BLACK
## 28 55 d BLACK
## 29 55 e BLACK
## 30 55 f BLACK
## 31 63 a BLACK
## 32 63 b BLUE
## 33 63 c BLUE
## 34 63 d BLUE
## 35 63 e BLUE
## 36 63 f BLUE
```

Data Exploration

```
str(df)
```

```
## 'data.frame': 36 obs. of 3 variables:
## $ X : Factor w/ 6 levels "5","19","35",...: 1 1 1 1 1 1 2 2 2 2 ...
## $ Y : Factor w/ 6 levels "a","b","c","d",...: 1 2 3 4 5 6 1 2 3 4 ...
## $ label: Factor w/ 2 levels "BLACK","BLUE": 2 1 2 1 1 1 2 2 2 2 ...
```

```
skim(df)
```

Table 1: Data summary

Name	df
Number of rows	36
Number of columns	3
Column type frequency:	
factor	3
Group variables	None

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
X	0	1	FALSE	6	5: 6, 19: 6, 35: 6, 51: 6
Y	0	1	FALSE	6	a: 6, b: 6, c: 6, d: 6
label	0	1	FALSE	2	BLA: 22, BLU: 14

```
summary(df)
```

```
##      X      Y      label
##  5 :6    a:6  BLACK:22
## 19:6    b:6  BLUE :14
## 35:6    c:6
## 51:6    d:6
## 55:6    e:6
## 63:6    f:6
```

Preparing the data for ML model

```
library(caret)
library(ModelMetrics)

respCol <- ncol(df)[[1]]
train <- createDataPartition(df[,respCol], p = .70) # training data
obs <- df[-train$Resample1, respCol] # test data

perfALG <- c("LR", "NB", "Knn3", "knn5")
perfAUC = numeric()
perfACC = numeric()
perfTPR = numeric()
perfFPR = numeric()
perfTNR = numeric()
perfFNR = numeric()
```

Simple Linear regression Model

```
## Logistic Regression (LR) Model
lrFit <- glm(label~., data = df[train$Resample1, ], family = binomial)
summary(lrFit)

##
## Call:
## glm(formula = label ~ ., family = binomial, data = df[train$Resample1,
##      ])
##
## Deviance Residuals:
```

```
##      Min      1Q      Median      3Q      Max
## -1.17741 -0.00003  0.00000  0.00003  1.17741
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -4.274e+01  1.881e+04 -0.002   0.998
## X19          4.274e+01  1.881e+04  0.002   0.998
## X35          9.642e-01  2.236e+04  0.000   1.000
## X51          8.598e-02  2.729e+04  0.000   1.000
## X55         -4.146e+01  2.615e+04 -0.002   0.999
## X63          4.274e+01  1.881e+04  0.002   0.998
## Yb           2.186e+01  1.901e+04  0.001   0.999
## Yc           6.340e+01  2.315e+04  0.003   0.998
## Yd           2.098e+01  1.369e+04  0.002   0.999
## Ye           1.583e-15  2.000e+00  0.000   1.000
## Yf           2.109e+01  1.820e+04  0.001   0.999
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 34.6465  on 25  degrees of freedom
## Residual deviance:  5.5452  on 15  degrees of freedom
## AIC: 27.545
##
## Number of Fisher Scoring iterations: 21
```

Fit the data into LR

```
lrProb <- predict(lrFit, newdata = df[-train$Resample1,], type = "response")

lrPred <- rep("BLACK", length(lrProb))
lrPred[lrProb > 0.5] = "BLUE"
lrPred = as.factor(lrPred)
#postResample(lrPred, obs)
perfAUC <- c(perfAUC, auc(actual = obs, predicted = lrPred))
perfACC <- c(perfACC, postResample(lrPred, obs)["Accuracy"])
perfTPR <- c(perfTPR, caret::sensitivity(lrPred, obs))
perfFPR <- c(perfFPR, 1 - caret::specificity(lrPred, obs))
perfTNR <- c(perfTNR, caret::specificity(lrPred, obs))
perfFNR <- c(perfFNR, 1 - caret::sensitivity(lrPred, obs))
table(lrPred, obs)
```

```
##      obs
## lrPred BLACK BLUE
##  BLACK     6    1
##  BLUE     0    3
```

As we can see, we got a confusion matrix with 6 True BLACKS, 3 True BLUES. However, we got 1 False BLUE predicted as BLACK. Over all the performance is not bad and considered a good classifier.

Naive Bayes Model

```
## Naive Bayes (NB) Model
library(e1071)
nbFit <- naiveBayes(label ~ ., data = df[train$Resample1, ])
#print(nbFit)
nbPred <- predict(nbFit, newdata = df[-train$Resample1,], type = "class")
#postResample(nbPred, obs)
perfAUC <- c(perfAUC, auc(actual = obs, predicted = nbPred))
perfACC <- c(perfACC, postResample(nbPred, obs)["Accuracy"])
perfTPR <- c(perfTPR, caret::sensitivity(nbPred, obs))
perfFPR <- c(perfFPR, 1 - caret::specificity(nbPred, obs))
perfTNR <- c(perfTNR, caret::specificity(nbPred, obs))
perfFNR <- c(perfFNR, 1 - caret::sensitivity(nbPred, obs))
table(nbPred, obs)
```

```
##      obs
## nbPred BLACK BLUE
##  BLACK      6    2
##  BLUE       0    2
```

It surprisingly introduces a higher prediction errors; although NB performs well with text classification and with a small data. It gave 6 Trues on BLACKS, 2 Trues on BLUE. However, it gave 2 Falses on BLUES predicted as BLACK.

KNN ML Model(k=3)

```
## KNN Model
knnFit3 <- knn3(label ~., data = df[train$Resample1,], k = 3)
knnPred3 <- predict(knnFit3, newdata = df[-train$Resample1,], type = "class")
#postResample(knnPred3, obs)
perfAUC <- c(perfAUC, auc(actual = obs, predicted = knnPred3))
perfACC <- c(perfACC, postResample(knnPred3, obs)["Accuracy"])
perfTPR <- c(perfTPR, caret::sensitivity(knnPred3, obs))
perfFPR <- c(perfFPR, 1 - caret::specificity(knnPred3, obs))
perfTNR <- c(perfTNR, caret::specificity(knnPred3, obs))
perfFNR <- c(perfFNR, 1 - caret::sensitivity(knnPred3, obs))
table(knnPred3, obs)
```

```
##      obs
## knnPred3 BLACK BLUE
##  BLACK      6    2
##  BLUE       0    2
```

KNN with K = 3 gave the same performance as Naive Bayes.

KNN ML Model (k=5)

```
## KNN Model
knnFit5 <- knn3(label ~., data = df[train$Resample1,], k = 5)
knnPred5 <- predict(knnFit5, newdata = df[-train$Resample1,], type = "class")
#postResample(knnPred, obs)
perfAUC <- c(perfAUC, auc(actual = obs, predicted = knnPred5))
perfACC <- c(perfACC, postResample(knnPred5, obs)["Accuracy"])
perfTPR <- c(perfTPR, caret::sensitivity(knnPred5, obs))
perfFPR <- c(perfFPR, 1 - caret::specificity(knnPred5, obs))
perfTNR <- c(perfTNR, caret::specificity(knnPred5, obs))
perfFNR <- c(perfFNR, 1 - caret::sensitivity(knnPred5, obs))
table(knnPred5, obs)
```

```
##      obs
## knnPred5 BLACK BLUE
##      BLACK      6    2
##      BLUE      0    2
```

Again, changing the k parameter didn't change much in the overall performance to the classifier, where it gave the same k = 3

Generate an accuracy comparing table

```
perf <- data.frame(
  ALGO = perfALG,
  AUC = perfAUC,
  ACCURACY = perfACC,
  TPR = perfTPR,
  FPR = perfFPR,
  TNR = perfTNR,
  FNR = perfFNR
)
```

```
perf
```

```
##   ALGO   AUC ACCURACY TPR  FPR  TNR  FNR
## 1  LR 0.875      0.9   1 0.25 0.75   0
## 2  NB 0.750      0.8   1 0.50 0.50   0
## 3 Knn3 0.750      0.8   1 0.50 0.50   0
## 4 knn5 0.750      0.8   1 0.50 0.50   0
```

Summary

It is hard to determine a clear winner among the classifiers, as multiple runs will select different training and testing data and greatly influence the training and performance of each model from run to run. However, since the dataset is pretty small and small datasets require strong assumption (bias). As it is preferable to use Occam's razor first. The less the assumptions are and the hypothesis is, the better is the results. In our case, Linear regression seems to perform well based on the AUC and ACCURACY.